

# Grouping and Summarising Dataframes

Grouping and aggregation are some of the most frequently used operations in data analysis, especially while doing exploratory data analysis (EDA), where comparing summary statistics across groups of data is common.

For e.g., in the retail sales data we are working with, you may want to compare the average sales of various regions, or compare the total profit of two customer segments.

Grouping analysis can be thought of as having three parts:

1. **Splitting** the data into groups (e.g. groups of customer segments, product categories, etc.)
2. **Applying** a function to each group (e.g. mean or total sales of each customer segment)
3. **Combining** the results into a data structure showing the summary statistics

Let's work through some examples.

```
In [1]: # Loading libraries and files
import numpy as np
import pandas as pd

market_df = pd.read_csv("../global_sales_data/market_fact.csv")
customer_df = pd.read_csv("../global_sales_data/cust_dimen.csv")
product_df = pd.read_csv("../global_sales_data/prod_dimen.csv")
shipping_df = pd.read_csv("../global_sales_data/shipping_dimen.csv")
orders_df = pd.read_csv("../global_sales_data/orders_dimen.csv")
```

Say you want to understand how well or poorly the business is doing in various customer segments, regions, product categories etc. Specifically, you want to identify areas of business where you are incurring heavy losses, and want to take action accordingly.

To do that, we will answer questions such as:

- Which customer segments are the least profitable?
- Which product categories and sub-categories are the least profitable?
- Customers in which geographic region cause the most losses?
- Etc.

First, we will merge all the dataframes, so we have all the data in one master\_df.

```
In [2]: # Merging the dataframes one by one
df_1 = pd.merge(market_df, customer_df, how='inner', on='Cust_id')
df_2 = pd.merge(df_1, product_df, how='inner', on='Prod_id')
df_3 = pd.merge(df_2, shipping_df, how='inner', on='Ship_id')
master_df = pd.merge(df_3, orders_df, how='inner', on='Ord_id')

master_df.head()
```

Out[2]:

	Ord_id	Prod_id	Ship_id	Cust_id	Sales	Discount	Order_Quantity	Profit
0	Ord_5446	Prod_16	SHP_7609	Cust_1818	136.81	0.01	23	-30.51
1	Ord_5446	Prod_4	SHP_7610	Cust_1818	4701.69	0.00	26	1148.90
2	Ord_5446	Prod_6	SHP_7608	Cust_1818	164.02	0.03	23	-47.64
3	Ord_2978	Prod_16	SHP_4112	Cust_1088	305.05	0.04	27	23.12
4	Ord_5484	Prod_16	SHP_7663	Cust_1820	322.82	0.05	35	-17.58

5 rows × 22 columns

### Step 1. Grouping using df.groupby()

Typically, you group the data using a categorical variable, such as customer segments, product categories, etc. This creates as many subsets of the data as there are levels in the categorical variable.

For example, in this case, we will group the data along Customer\_Segment.

```
In [3]: # Which customer segments are the least profitable?

# Step 1. Grouping: First, we will group the dataframe by customer segments
df_by_segment = master_df.groupby('Customer_Segment')
df_by_segment
```

Out[3]: <pandas.core.groupby.DataFrameGroupBy object at 0x1046be710>

Note that df.groupby returns a DataFrameGroupBy object.

### Step 2. Applying a Function

After grouping, you apply a function to a **numeric variable**, such as mean(Sales), sum(Profit), etc.

```
In [4]: # Step 2. Applying a function  
# We can choose aggregate functions such as sum, mean, median, etc.  
df_by_segment['Profit'].sum()
```

```
Out[4]: Customer_Segment  
CONSUMER          287959.94  
CORPORATE          599746.00  
HOME OFFICE        318354.03  
SMALL BUSINESS     315708.01  
Name: Profit, dtype: float64
```

Notice that we have indexed the Profit column in the DataFrameGroupBy object exactly as we index a normal column in a dataframe. Alternatively, you could also use `df_by_segment.Profit`.

```
In [5]: # Alternatively  
df_by_segment.Profit.sum()
```

```
Out[5]: Customer_Segment  
CONSUMER          287959.94  
CORPORATE          599746.00  
HOME OFFICE        318354.03  
SMALL BUSINESS     315708.01  
Name: Profit, dtype: float64
```

So this tells us that profits are the least in the CONSUMER segment, and highest in the CORPORATE segment.

```
In [6]: # For better readability, you may want to sort the summarised series:  
df_by_segment.Profit.sum().sort_values(ascending = False)
```

```
Out[6]: Customer_Segment  
CORPORATE          599746.00  
HOME OFFICE        318354.03  
SMALL BUSINESS     315708.01  
CONSUMER          287959.94  
Name: Profit, dtype: float64
```

### Step 3. Combining the results into a Data Structure

You can optionally show the results as a dataframe.

```
In [7]: # Converting to a df  
pd.DataFrame(df_by_segment['Profit'].sum())
```

Out[7]:

	Profit
Customer_Segment	
CONSUMER	287959.94
CORPORATE	599746.00
HOME OFFICE	318354.03
SMALL BUSINESS	315708.01

```
In [8]: # Let's go through some more examples  
# E.g.: Which product categories are the least profitable?  
  
# 1. Group by product category  
by_product_cat = master_df.groupby('Product_Category')
```

```
In [9]: # 2. This time, let's compare average profits  
# Apply mean() on Profit  
by_product_cat['Profit'].mean()
```

```
Out[9]: Product_Category  
FURNITURE      68.116607  
OFFICE SUPPLIES 112.369074  
TECHNOLOGY     429.207516  
Name: Profit, dtype: float64
```

FURNITURE is the least profitable, TECHNOLOGY the most. Let's see which product sub-categories within FURNITURE are less profitable.

```
In [10]: # E.g.: Which product categories and sub-categories are the least profitable?
# 1. Group by category and sub-category
by_product_cat_subcat = master_df.groupby(['Product_Category', 'Product_Sub_Category'])
by_product_cat_subcat['Profit'].mean()
```

```
Out[10]: Product_Category  Product_Sub_Category
FURNITURE                BOOKCASES                -177.683228
                CHAIRS & CHAIRMATS                387.693601
                OFFICE FURNISHINGS                127.446612
                TABLES                -274.411357
OFFICE SUPPLIES          APPLIANCES                223.866498
                BINDERS AND BINDER ACCESSORIES        335.970918
                ENVELOPES                195.864228
                LABELS                47.490174
                PAPER                36.949551
                PENS & ART SUPPLIES                11.950679
                RUBBER BANDS                -0.573575
                SCISSORS, RULERS AND TRIMMERS        -54.161458
                STORAGE & ORGANIZATION                12.205403
TECHNOLOGY              COMPUTER PERIPHERALS        124.389815
                COPIERS AND FAX                1923.695287
                OFFICE MACHINES                913.094748
                TELEPHONES AND COMMUNICATION        358.948607
Name: Profit, dtype: float64
```

Thus, within FURNITURE, TABLES are the least profitable, followed by BOOKCASES.

```
In [11]: # Recall the df.describe() method?
# To apply multiple functions simultaneously, you can use the describe() function on the grouped df object
by_product_cat['Profit'].describe()
```

```
Out[11]:
```

	count	mean	std	min	25%	50%	7
<b>Product_Category</b>							
<b>FURNITURE</b>	1724.0	68.116607	1112.923257	-11053.60	-281.3550	-14.250	187.16
<b>OFFICE SUPPLIES</b>	4610.0	112.369074	744.617939	-2175.09	-57.0225	-3.845	56.947
<b>TECHNOLOGY</b>	2065.0	429.207516	1863.208375	-14140.70	-88.9400	66.220	561.13

```
In [12]: # Some other summary functions to apply on groups
by_product_cat['Profit'].count()
```

```
Out[12]: Product_Category
FURNITURE                1724
OFFICE SUPPLIES          4610
TECHNOLOGY               2065
Name: Profit, dtype: int64
```

```
In [13]: by_product_cat['Profit'].min()
```

```
Out[13]: Product_Category
FURNITURE      -11053.60
OFFICE SUPPLIES -2175.09
TECHNOLOGY     -14140.70
Name: Profit, dtype: float64
```

```
In [14]: # E.g. Customers in which geographic region are the least profitable?
master_df.groupby('Region').Profit.mean()
```

```
Out[14]: Region
ATLANTIC      221.259870
NORTHWEST TERRITORIES 255.464670
NUNAVUT       35.963418
ONTARIO      189.960865
PRARIE       188.253294
QUEBEC       179.803649
WEST        149.175595
YUKON       136.253155
Name: Profit, dtype: float64
```

```
In [15]: # Note that the resulting object is a Series, thus you can perform vectorised
          computations on them
```

```
# E.g. Calculate the Sales across each region as a percentage of total Sales
# You can divide the entire series by a number (total sales) easily
(master_df.groupby('Region').Sales.sum() / sum(master_df['Sales']))*100
```

```
Out[15]: Region
ATLANTIC      13.504305
NORTHWEST TERRITORIES  5.369193
NUNAVUT       0.780233
ONTARIO      20.536970
PRARIE       19.022396
QUEBEC       10.124936
WEST        24.119372
YUKON        6.542595
Name: Sales, dtype: float64
```

The regions ONTARIO, WEST and PRARIE comprise of about 64% of the sales.

Until now, we've been working with the data without making changes or additions to it. In the next section, we will create new columns, alter existing columns and apply some more grouping and summarising.