

Car Price Prediction - Assignment Solution ¶

The solution is divided into the following sections:

- Data understanding and exploration
- Data cleaning
- Data preparation
- Model building and evaluation

1. Data Understanding and Exploration

Let's first have a look at the dataset and understand the size, attribute names etc.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV

import os

# hide warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # reading the dataset
cars = pd.read_csv("CarPrice_Assignment.csv")
```

In [3]: *# summary of the dataset: 205 rows, 26 columns, no null values*

```
print(cars.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
car_ID                205 non-null int64
symboling             205 non-null int64
CarName               205 non-null object
fueltype              205 non-null object
aspiration             205 non-null object
doornumber            205 non-null object
carbody               205 non-null object
drivewheel            205 non-null object
enginelocation        205 non-null object
wheelbase             205 non-null float64
carlength             205 non-null float64
carwidth              205 non-null float64
carheight             205 non-null float64
curbweight            205 non-null int64
enginetype            205 non-null object
cylindernumber        205 non-null object
enginesize            205 non-null int64
fuelsystem            205 non-null object
boreratio             205 non-null float64
stroke                205 non-null float64
compressionratio      205 non-null float64
horsepower            205 non-null int64
peakrpm               205 non-null int64
citympg               205 non-null int64
highwaympg            205 non-null int64
price                 205 non-null float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.7+ KB
None
```

```
In [4]: # head
cars.head()
```

```
Out[4]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewhe
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd
3	4	2	audi 100 ls	gas	std	four	sedan	fwd
4	5	2	audi 100ls	gas	std	four	sedan	4wd

5 rows × 26 columns

Understanding the Data Dictionary

The data dictionary contains the meaning of various attributes; some non-obvious ones are:

```
In [5]: # symboling: -2 (least risky) to +3 most risky
# Most cars are 0,1,2
cars['symboling'].astype('category').value_counts()
```

```
Out[5]: 0    67
        1    54
        2    32
        3    27
       -1    22
       -2     3
Name: symboling, dtype: int64
```

```
In [6]: # aspiration: An (internal combustion) engine property showing
# whether the oxygen intake is through standard (atmospheric pressure)
# or through turbocharging (pressurised oxygen intake)


cars['aspiration'].astype('category').value_counts()
```

```
Out[6]: std    168
        turbo   37
Name: aspiration, dtype: int64
```


```
In [7]: # drivewheel: frontwheel, rarewheel or four-wheel drive  
cars['drivewheel'].astype('category').value_counts()
```

```
Out[7]: fwd      120  
       rwd       76  
       4wd        9  
       Name: drivewheel, dtype: int64
```


```
In [8]: # wheelbase: distance between centre of front and rarewheels  
sns.distplot(cars['wheelbase'])  
plt.show()
```




```
In [9]: # curbweight: weight of car without occupants or baggage  
sns.distplot(cars['curbweight'])  
plt.show()
```




```
In [10]: # stroke: volume of the engine (the distance traveled by the  
# piston in each cycle)  
sns.distplot(cars['stroke'])  
plt.show()
```



```
In [11]: # compression ratio: ratio of volume of compression chamber  
# at largest capacity to least capacity  
sns.distplot(cars['compressionratio'])  
plt.show()
```



```
In [12]: # target variable: price of car
sns.distplot(cars['price'])
plt.show()
```



Data Exploration

To perform linear regression, the (numeric) target variable should be linearly related to *at least one another numeric variable*. Let's see whether that's true in this case.

We'll first subset the list of all (independent) numeric variables, and then make a **pairwise plot**.

```
In [20]: # all numeric (float and int) variables in the dataset
cars_numeric = cars.select_dtypes(include=['float64', 'int64'])
cars_numeric.head()
```

Out[20]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize
0	1	3	88.6	168.8	64.1	48.8	2548	130
1	2	3	88.6	168.8	64.1	48.8	2548	130
2	3	1	94.5	171.2	65.5	52.4	2823	152
3	4	2	99.8	176.6	66.2	54.3	2337	109
4	5	2	99.4	176.6	66.4	54.3	2824	136

Here, although the variable `symboling` is numeric (int), we'd rather treat it as categorical since it has only 6 discrete values. Also, we do not want 'car_ID'.

```
In [21]: # dropping symboling and car_ID
cars_numeric = cars_numeric.drop(['symboling'], axis=1)
cars_numeric.head()
```

Out[21]:

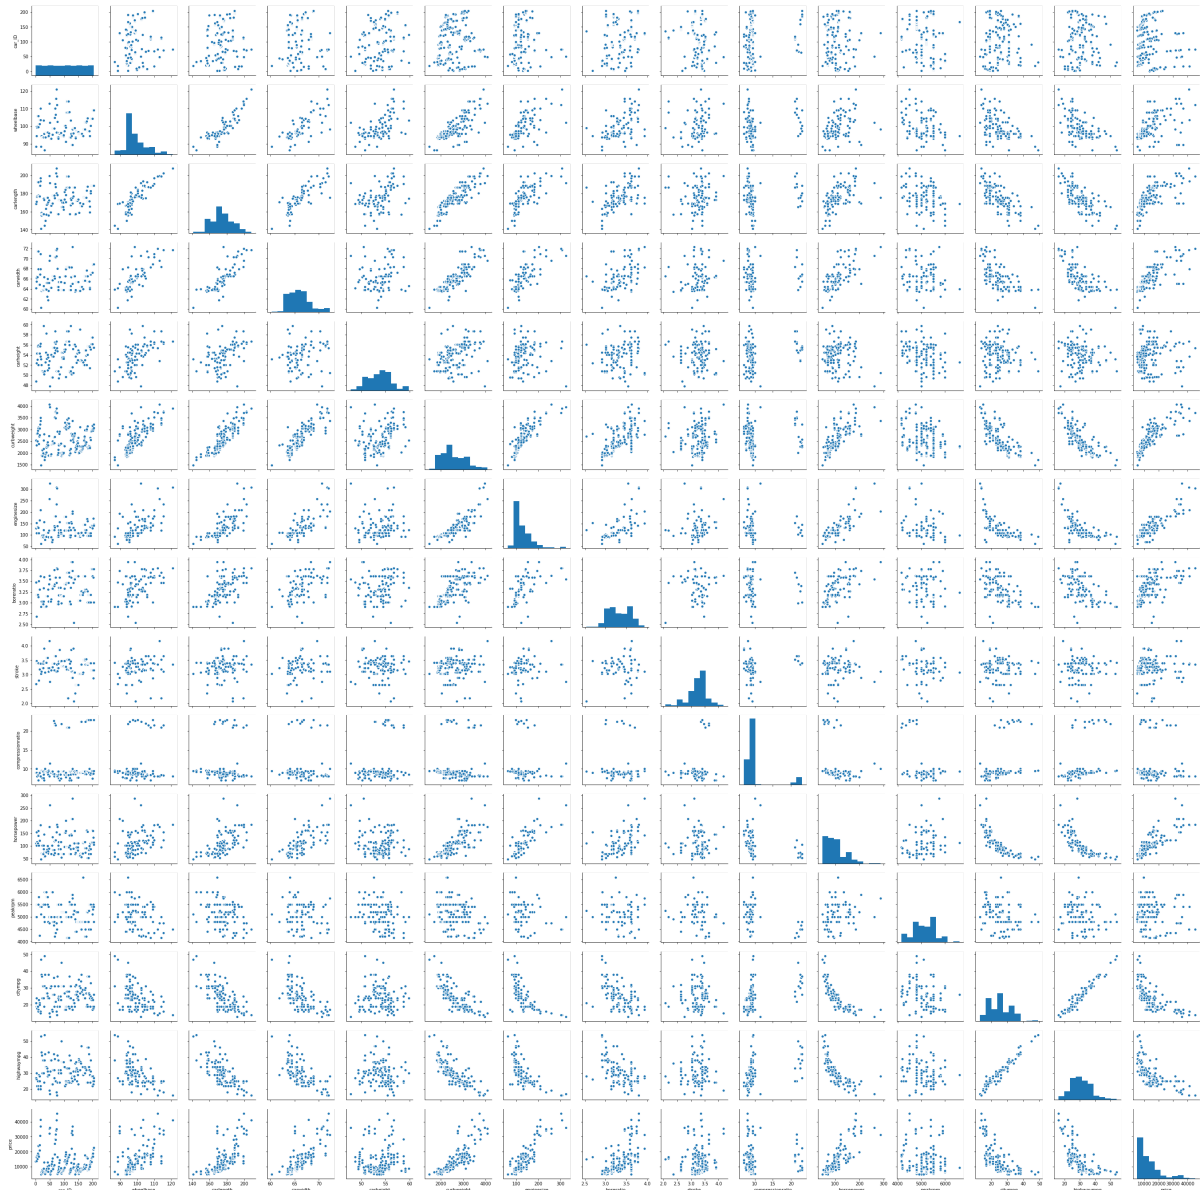
	car_ID	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	
0	1	88.6	168.8	64.1	48.8	2548	130	3.47	:
1	2	88.6	168.8	64.1	48.8	2548	130	3.47	:
2	3	94.5	171.2	65.5	52.4	2823	152	2.68	:
3	4	99.8	176.6	66.2	54.3	2337	109	3.19	:
4	5	99.4	176.6	66.4	54.3	2824	136	3.19	:

Let's now make a pairwise scatter plot and observe linear relationships.

```
In [22]: # pairwise scatter plot

plt.figure(figsize=(20, 10))
sns.pairplot(cars_numeric)
plt.show()
```

<matplotlib.figure.Figure at 0x24e9998a278>



This is quite hard to read, and we can rather plot correlations between variables. Also, a heatmap is pretty useful to visualise multiple correlations in one plot.


```
In [23]: # correlation matrix
cor = cars_numeric.corr()
cor
```


Out[23]:

	car_ID	wheelbase	carlength	carwidth	carheight	curbweight	engine
car_ID	1.000000	0.129729	0.170636	0.052387	0.255960	0.071962	-0.000000
wheelbase	0.129729	1.000000	0.874587	0.795144	0.589435	0.776386	0.500000
carlength	0.170636	0.874587	1.000000	0.841118	0.491029	0.877728	0.600000
carwidth	0.052387	0.795144	0.841118	1.000000	0.279210	0.867032	0.700000
carheight	0.255960	0.589435	0.491029	0.279210	1.000000	0.295572	0.000000
curbweight	0.071962	0.776386	0.877728	0.867032	0.295572	1.000000	0.800000
engine	-0.033930	0.569329	0.683360	0.735433	0.067149	0.850594	1.000000
bore	0.260064	0.488750	0.606454	0.559150	0.171071	0.648480	0.500000
stroke	-0.160824	0.160959	0.129533	0.182942	-0.055307	0.168790	0.200000
compressionratio	0.150276	0.249786	0.158414	0.181129	0.261214	0.151362	0.000000
horsepower	-0.015006	0.353294	0.552623	0.640732	-0.108802	0.750739	0.800000
peakrpm	-0.203789	-0.360469	-0.287242	-0.220012	-0.320411	-0.266243	-0.000000
citympg	0.015940	-0.470414	-0.670909	-0.642704	-0.048640	-0.757414	-0.000000
highwaympg	0.011255	-0.544082	-0.704662	-0.677218	-0.107358	-0.797465	-0.000000
price	-0.109093	0.577816	0.682920	0.759325	0.119336	0.835305	0.800000

```
In [24]: # plotting correlations on a heatmap

# figure size
plt.figure(figsize=(16,8))

# heatmap
sns.heatmap(cor, cmap="YlGnBu", annot=True)
plt.show()
```



The heatmap shows some useful insights:

Correlation of price with independent variables:

- Price is highly (positively) correlated with wheelbase, carlength, carwidth, curbweight, enginesize, horsepower (notice how all of these variables represent the size/weight/engine power of the car)
- Price is negatively correlated to citympg and highwaympg (-0.70 approximately). This suggest that cars having high mileage may fall in the 'economy' cars category, and are priced lower (think Maruti Alto/Swift type of cars, which are designed to be affordable by the middle class, who value mileage more than horsepower/size of car etc.)

Correlation among independent variables:

- Many independent variables are highly correlated (look at the top-left part of matrix): wheelbase, carlength, curbweight, enginesize etc. are all measures of 'size/weight', and are positively correlated

Thus, while building the model, we'll have to pay attention to multicollinearity (especially linear models, such as linear and logistic regression, suffer more from multicollinearity).

2. Data Cleaning

Let's now conduct some data cleaning steps.

We've seen that there are no missing values in the dataset. We've also seen that variables are in the correct format, except symboling, which should rather be a categorical variable (so that dummy variable are created for the categories).

Note that it *can* be used in the model as a numeric variable also.

```
In [25]: # variable formats
cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
car_ID          205 non-null int64
symboling       205 non-null int64
CarName         205 non-null object
fueltype        205 non-null object
aspiration      205 non-null object
doornumber      205 non-null object
carbody         205 non-null object
drivewheel      205 non-null object
enginelocation  205 non-null object
wheelbase       205 non-null float64
carlength       205 non-null float64
carwidth        205 non-null float64
carheight       205 non-null float64
curbweight      205 non-null int64
enginetype      205 non-null object
cylindernumber  205 non-null object
enginesize      205 non-null int64
fuelsystem      205 non-null object
boreratio       205 non-null float64
stroke          205 non-null float64
compressionratio 205 non-null float64
horsepower      205 non-null int64
peakrpm         205 non-null int64
citympg         205 non-null int64
highwaympg      205 non-null int64
price           205 non-null float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.7+ KB
```

```
In [26]: # converting symboling to categorical
cars['symboling'] = cars['symboling'].astype('object')
cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
car_ID                205 non-null int64
symboling             205 non-null object
CarName               205 non-null object
fueltype              205 non-null object
aspiration            205 non-null object
doornumber            205 non-null object
carbody               205 non-null object
drivewheel            205 non-null object
enginelocation        205 non-null object
wheelbase             205 non-null float64
carlength             205 non-null float64
carwidth              205 non-null float64
carheight             205 non-null float64
curbweight            205 non-null int64
engine               205 non-null object
cylindernumber        205 non-null object
enginesize            205 non-null int64
fuelsystem            205 non-null object
boreratio             205 non-null float64
stroke                205 non-null float64
compressionratio      205 non-null float64
horsepower            205 non-null int64
peakrpm               205 non-null int64
citympg               205 non-null int64
highwaympg            205 non-null int64
price                 205 non-null float64
dtypes: float64(8), int64(7), object(11)
memory usage: 41.7+ KB
```

Next, we need to extract the company name from the column CarName.

```
In [27]: # CarName: first few entries  
cars['CarName'][:30]
```

```
Out[27]: 0      alfa-romero giulia  
1      alfa-romero stelvio  
2      alfa-romero Quadrifoglio  
3      audi 100 ls  
4      audi 100ls  
5      audi fox  
6      audi 100ls  
7      audi 5000  
8      audi 4000  
9      audi 5000s (diesel)  
10     bmw 320i  
11     bmw 320i  
12     bmw x1  
13     bmw x3  
14     bmw z4  
15     bmw x4  
16     bmw x5  
17     bmw x3  
18     chevrolet impala  
19     chevrolet monte carlo  
20     chevrolet vega 2300  
21     dodge rampage  
22     dodge challenger se  
23     dodge d200  
24     dodge monaco (sw)  
25     dodge colt hardtop  
26     dodge colt (sw)  
27     dodge coronet custom  
28     dodge dart custom  
29     dodge coronet custom (sw)  
Name: CarName, dtype: object
```

Notice that the carname is what occurs before a space, e.g. alfa-romero, audi, chevrolet, dodge, bmx etc.

Thus, we need to simply extract the string before a space. There are multiple ways to do that.

```
In [28]: # Extracting carname

# Method 1: str.split() by space
carnames = cars['CarName'].apply(lambda x: x.split(" ")[0])
carnames[:30]
```

```
Out[28]: 0    alfa-romero
1    alfa-romero
2    alfa-romero
3         audi
4         audi
5         audi
6         audi
7         audi
8         audi
9         audi
10        bmw
11        bmw
12        bmw
13        bmw
14        bmw
15        bmw
16        bmw
17        bmw
18    chevrolet
19    chevrolet
20    chevrolet
21        dodge
22        dodge
23        dodge
24        dodge
25        dodge
26        dodge
27        dodge
28        dodge
29        dodge
Name: CarName, dtype: object
```

```
In [29]: # Method 2: Use regular expressions  
import re  
  
# regex: any alphanumeric sequence before a space, may contain a hyphen  
p = re.compile(r'\w+~?\w+')  
carnames = cars['CarName'].apply(lambda x: re.findall(p, x)[0])  
print(carnames)
```

0	alfa-romero
1	alfa-romero
2	alfa-romero
3	audi
4	audi
5	audi
6	audi
7	audi
8	audi
9	audi
10	bmw
11	bmw
12	bmw
13	bmw
14	bmw
15	bmw
16	bmw
17	bmw
18	chevrolet
19	chevrolet
20	chevrolet
21	dodge
22	dodge
23	dodge
24	dodge
25	dodge
26	dodge
27	dodge
28	dodge
29	dodge
	...
175	toyota
176	toyota
177	toyota
178	toyota
179	toyota
180	toyota
181	toyouta
182	vokswagen
183	volkswagen
184	volkswagen
185	volkswagen
186	volkswagen
187	volkswagen
188	volkswagen
189	vw
190	vw
191	volkswagen
192	volkswagen
193	volkswagen
194	volvo
195	volvo
196	volvo
197	volvo
198	volvo
199	volvo
200	volvo


```

201      volvo
202      volvo
203      volvo
204      volvo
Name: CarName, Length: 205, dtype: object

```

Let's create a new column to store the company name and check whether it looks okay.

```

In [30]: # New column car_company
cars['car_company'] = cars['CarName'].apply(lambda x: re.findall(p, x)[0])

```

```

In [31]: # Look at all values
cars['car_company'].astype('category').value_counts()

```

```

Out[31]: toyota      31
         nissan      17
         mazda      15
         mitsubishi  13
         honda      13
         subaru     12
         volvo      11
         peugeot    11
         volkswagen  9
         dodge      9
         buick      8
         bmw        8
         plymouth   7
         audi       7
         saab       6
         porsche    4
         isuzu      4
         chevrolet  3
         alfa-romero 3
         jaguar     3
         vw        2
         maxda     2
         renault   2
         mercury   1
         porcshee  1
         toyouta   1
         vokswagen 1
         Nissan    1
Name: car_company, dtype: int64

```

Notice that **some car-company names are misspelled** - vw and vokswagen should be volkswagen, porcshee should be porsche, toyouta should be toyota, Nissan should be nissan, maxda should be mazda etc.

This is a data quality issue, let's solve it.

```
In [32]: # replacing misspelled car_company names

# volkswagen
cars.loc[(cars['car_company'] == "vw") |
         (cars['car_company'] == "vokswagen"), 'car_company'] = 'volkswagen'

# porsche
cars.loc[cars['car_company'] == "porcshce", 'car_company'] = 'porsche'

# toyota
cars.loc[cars['car_company'] == "toyouta", 'car_company'] = 'toyota'

# nissan
cars.loc[cars['car_company'] == "Nissan", 'car_company'] = 'nissan'

# mazda
cars.loc[cars['car_company'] == "maxda", 'car_company'] = 'mazda'
```

```
In [33]: cars['car_company'].astype('category').value_counts()
```

```
Out[33]: toyota      32
         nissan      18
         mazda      17
         honda      13
         mitsubishi  13
         subaru      12
         volkswagen  12
         volvo       11
         peugeot     11
         dodge        9
         buick        8
         bmw          8
         plymouth     7
         audi         7
         saab         6
         porsche      5
         isuzu        4
         alfa-romero  3
         chevrolet    3
         jaguar       3
         renault      2
         mercury      1
         Name: car_company, dtype: int64
```

The car_company variable looks okay now. Let's now drop the car name variable.

```
In [34]: # drop carname variable
cars = cars.drop('CarName', axis=1)
```

In [35]: cars.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
car_ID          205 non-null int64
symboling       205 non-null object
fueltype        205 non-null object
aspiration      205 non-null object
doornumber      205 non-null object
carbody         205 non-null object
drivewheel      205 non-null object
enginelocation  205 non-null object
wheelbase       205 non-null float64
carlength       205 non-null float64
carwidth        205 non-null float64
carheight       205 non-null float64
curbweight      205 non-null int64
enginetype      205 non-null object
cylindernumber  205 non-null object
enginesize      205 non-null int64
fuelsystem      205 non-null object
boreratio       205 non-null float64
stroke          205 non-null float64
compressionratio 205 non-null float64
horsepower      205 non-null int64
peakrpm         205 non-null int64
citympg         205 non-null int64
highwaympg      205 non-null int64
price           205 non-null float64
car_company     205 non-null object
dtypes: float64(8), int64(7), object(11)
memory usage: 41.7+ KB
```

In [36]: # outliers
cars.describe()

Out[36]:

	car_ID	wheelbase	carlength	carwidth	carheight	curbweight	engine
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	98.756585	174.049268	65.907805	53.724878	2555.565854	126.900000
std	59.322565	6.021776	12.337289	2.145204	2.443522	520.680204	41.642000
min	1.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000
25%	52.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000
50%	103.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000
75%	154.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000
max	205.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000

```
In [37]: cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
car_ID                205 non-null int64
symboling             205 non-null object
fueltype              205 non-null object
aspiration            205 non-null object
doornumber            205 non-null object
carbody               205 non-null object
drivewheel            205 non-null object
engineloation         205 non-null object
wheelbase             205 non-null float64
carlength             205 non-null float64
carwidth              205 non-null float64
carheight             205 non-null float64
curbweight            205 non-null int64
enginetype            205 non-null object
cylindernumber        205 non-null object
enginesize            205 non-null int64
fuelsystem            205 non-null object
boreratio             205 non-null float64
stroke                205 non-null float64
compressionratio      205 non-null float64
horsepower            205 non-null int64
peakrpm              205 non-null int64
citympg               205 non-null int64
highwaympg            205 non-null int64
price                 205 non-null float64
car_company           205 non-null object
dtypes: float64(8), int64(7), object(11)
memory usage: 41.7+ KB
```

3. Data Preparation

Data Preparation

Let's now prepare the data and build the model.

```
In [45]: # split into X and y
X = cars.loc[:, ['symboling', 'fueltype', 'aspiration', 'doornumber',
                 'carbody', 'drivewheel', 'engineloation', 'wheelbase', 'carlength',
                 'carwidth', 'carheight', 'curbweight', 'enginetype', 'cylindernumber',
                 'enginesize', 'fuelsystem', 'boreratio', 'stroke', 'compressionratio',
                 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
                 'car_company']]

y = cars['price']
```

In [46]: *# creating dummy variables for categorical variables*

```
# subset all categorical variables
cars_categorical = X.select_dtypes(include=['object'])
cars_categorical.head()
```

Out[46]:

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	engine
0	3	gas	std	two	convertible	rwd	front	cyl
1	3	gas	std	two	convertible	rwd	front	cyl
2	1	gas	std	two	hatchback	rwd	front	cyl
3	2	gas	std	four	sedan	fwd	front	cyl
4	2	gas	std	four	sedan	4wd	front	cyl

In [47]: *# convert into dummies*

```
cars_dummies = pd.get_dummies(cars_categorical, drop_first=True)
cars_dummies.head()
```

Out[47]:

	symboling_-1	symboling_0	symboling_1	symboling_2	symboling_3	fueltype_gas	fueltype_diesel
0	0	0	0	0	1	1	0
1	0	0	0	0	1	1	0
2	0	0	1	0	0	1	0
3	0	0	0	1	0	1	0
4	0	0	0	1	0	1	0

5 rows × 55 columns

In [48]: *# drop categorical variables*

```
X = X.drop(list(cars_categorical.columns), axis=1)
```

In [49]: *# concat dummy variables with X*

```
X = pd.concat([X, cars_dummies], axis=1)
```

```
In [50]: # scaling the features
from sklearn.preprocessing import scale

# storing column names in cols, since column names are (annoyingly) lost after
# scaling (the df is converted to a numpy array)
cols = X.columns
X = pd.DataFrame(scale(X))
X.columns = cols
X.columns
```

```
Out[50]: Index(['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight',
               'enginesize', 'boreratio', 'stroke', 'compressionratio', 'horsepower',
               'peakrpm', 'citympg', 'highwaympg', 'symboling_1', 'symboling_0',
               'symboling_1', 'symboling_2', 'symboling_3', 'fueltype_gas',
               'aspiration_turbo', 'doornumber_two', 'carbody_hardtop',
               'carbody_hatchback', 'carbody_sedan', 'carbody_wagon', 'drivewheel_fw
               d',
               'drivewheel_rwd', 'engine_location_rear', 'enginetype_dohcv',
               'enginetype_l', 'enginetype_ohc', 'enginetype_ohcf', 'enginetype_ohc
               v',
               'enginetype_rotor', 'cylindernumber_five', 'cylindernumber_four',
               'cylindernumber_six', 'cylindernumber_three', 'cylindernumber_twelve',
               'cylindernumber_two', 'fuelsystem_2bbl', 'fuelsystem_4bbl',
               'fuelsystem_idi', 'fuelsystem_mfi', 'fuelsystem_mphi',
               'fuelsystem_spdi', 'fuelsystem_spfi', 'car_company_audi',
               'car_company_bmw', 'car_company_buick', 'car_company_chevrolet',
               'car_company_dodge', 'car_company_honda', 'car_company_isuzu',
               'car_company_jaguar', 'car_company_mazda', 'car_company_mercury',
               'car_company_mitsubishi', 'car_company_nissan', 'car_company_peugeot',
               'car_company_plymouth', 'car_company_porsche', 'car_company_renault',
               'car_company_saab', 'car_company_subaru', 'car_company_toyota',
               'car_company_volkswagen', 'car_company_volvo'],
              dtype='object')
```

```
In [51]: # split into train and test
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    train_size=0.7,
                                                    test_size = 0.3, random_st
                                                    ate=100)
```

3. Model Building and Evaluation

Ridge and Lasso Regression

Let's now try predicting car prices, a dataset used in simple linear regression, to perform ridge and lasso regression.

Ridge Regression

```
In [52]: # list of alphas to tune
params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1,
                    0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
                    4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500, 1000 ]}

ridge = Ridge()

# cross validation
folds = 5
model_cv = GridSearchCV(estimator = ridge,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)
model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed: 0.4s finished

```
Out[52]: GridSearchCV(cv=5, error_score='raise',
                    estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=N
one,
                    normalize=False, random_state=None, solver='auto', tol=0.001),
                    fit_params=None, iid=True, n_jobs=1,
                    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4,
0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 2
0, 50, 100, 500, 1000]}},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='neg_mean_absolute_error', verbose=1)
```

```
In [53]: cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results = cv_results[cv_results['param_alpha']<=200]
cv_results.head()
```


Out[53]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_alpha
0	0.002807	0.000503	-2455.668504	-793.046184	0.0001
1	0.002207	0.000501	-2449.633986	-792.778547	0.001
2	0.003008	0.000702	-2404.435462	-790.876517	0.01
3	0.002013	0.000494	-2276.161219	-788.409662	0.05
4	0.001905	0.000401	-2190.708689	-787.299687	0.1

5 rows × 6 columns

```
In [54]: # plotting mean test and train scores with alpha
cv_results['param_alpha'] = cv_results['param_alpha'].astype('int32')

# plotting
plt.plot(cv_results['param_alpha'], cv_results['mean_train_score'])
plt.plot(cv_results['param_alpha'], cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.ylabel('Negative Mean Absolute Error')
plt.title("Negative Mean Absolute Error and alpha")
plt.legend(['train score', 'test score'], loc='upper left')
plt.show()
```




```
In [55]: alpha = 15
ridge = Ridge(alpha=alpha)

ridge.fit(X_train, y_train)
ridge.coef_
```

```
Out[55]: array([ 3.52557941e+02,  9.46530017e+01,  1.34511343e+03, -3.35849314e+02,
 1.16515775e+03,  1.33122540e+03, -1.18579444e+01, -2.84669666e+02,
 1.23313363e+01,  9.29481098e+02,  3.23313329e+02, -2.04711197e+01,
-8.41293665e+01,  2.21783691e+02,  1.66376475e+02,  6.11555386e+01,
-1.75468993e+01,  2.28698304e+02, -2.11820147e+02,  4.11053298e+02,
 1.01320701e+02,  2.63588369e+01, -5.29501748e+02, -2.71545584e+02,
-2.23762142e+02, -2.44396310e+02,  2.61656901e+02,  9.20191503e+02,
-1.93176012e+01, -2.18315631e+02,  3.13602886e+02,  2.38744211e+01,
 3.92105941e+00,  9.46640067e+01, -4.24237855e+02, -5.81409615e+02,
-1.05817250e+02,  2.93070489e+02, -2.36048487e+02,  9.46640067e+01,
 5.15255069e+01, -1.79077288e+02,  2.11820147e+02, -1.44440927e-28,
-3.95308641e+00, -1.36403005e+02, -1.44440927e-28,  3.35662090e+02,
 1.46679730e+03,  1.05750051e+03, -1.30931401e+02, -3.49739603e+02,
-2.82007527e+02, -3.98771447e+01,  8.28333109e+02, -2.30841923e+02,
-1.44440927e-28, -5.99411625e+02, -4.01098290e+02, -3.18047678e+02,
-2.57350351e+02,  6.20255051e+02, -2.19413859e+02,  1.89789295e+02,
-4.44216995e+02, -5.71822975e+02, -1.19897731e+02, -7.39919628e-01])
```

Lasso

```
In [56]: lasso = Lasso()

# cross validation
model_cv = GridSearchCV(estimator = lasso,
                        param_grid = params,
                        scoring= 'neg_mean_absolute_error',
                        cv = folds,
                        return_train_score=True,
                        verbose = 1)

model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

[Parallel(n_jobs=1)]: Done 140 out of 140 | elapsed: 1.9s finished

```
Out[56]: GridSearchCV(cv=5, error_score='raise',
                      estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1
000,
                      normalize=False, positive=False, precompute=False, random_state=None,
                      selection='cyclic', tol=0.0001, warm_start=False),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3, 0.4,
0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 2
0, 50, 100, 500, 1000]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_absolute_error', verbose=1)
```

```
In [57]: cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results.head()
```

Out[57]:


	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_alpha
0	0.020253	0.000703	-2200.571734	-792.317797	0.0001
1	0.015039	0.000402	-2200.636114	-792.315096	0.001
2	0.017645	0.000603	-2201.283963	-792.291655	0.01
3	0.014738	0.000502	-2204.263131	-792.251073	0.05
4	0.016644	0.000502	-2208.163365	-792.236231	0.1

5 rows × 6 columns

```
In [58]: # plotting mean test and train scores with alpha
cv_results['param_alpha'] = cv_results['param_alpha'].astype('float32')

# plotting
plt.plot(cv_results['param_alpha'], cv_results['mean_train_score'])
plt.plot(cv_results['param_alpha'], cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.ylabel('Negative Mean Absolute Error')

plt.title("Negative Mean Absolute Error and alpha")
plt.legend(['train score', 'test score'], loc='upper left')
plt.show()
```



```
In [59]: alpha =100

lasso = Lasso(alpha=alpha)

lasso.fit(X_train, y_train)
```

```
Out[59]: Lasso(alpha=100, copy_X=True, fit_intercept=True, max_iter=1000,
           normalize=False, positive=False, precompute=False, random_state=None,
           selection='cyclic', tol=0.0001, warm_start=False)
```

```
In [60]: lasso.coef_
```

```
Out[60]: array([  0.          , -0.          , 1747.1052243 , -82.23183774,
                1780.64173078,  788.28807799, -0.          , -0.          ,
                 0.          , 1017.48820119,  84.89633333,  0.          ,
                -0.          ,  0.          , -0.          , -0.          ,
                 0.          ,  246.519852 , -73.38572878, 120.56790634,
                 0.          ,  0.          , -187.60748943,  0.          ,
               -96.25412649, -134.39227325,  294.27227486, 1218.02281069,
                 0.          , -0.          ,  0.          , -0.          ,
                -0.          ,  0.          , -0.          , -202.47407284,
                -0.          , 197.70712322, -0.          ,  0.          ,
                -0.          , -0.          ,  58.81424436, -0.          ,
                 0.          , -0.          , -0.          , 186.35685239,
              1805.30123983, 1210.72936345,  0.          , -0.          ,
                -0.          ,  78.54297249,  796.29612837,  0.          ,
                -0.          , -397.80411254, -58.198149 , -377.78256238,
                -0.          ,  592.06274204, -163.73847377,  95.37139425,
              -198.09298955, -233.82794826,  0.          ,  206.40038676])
```